

A photograph of two dogs playing in a grassy field under a clear blue sky. On the left, a small brown and white dog is jumping or running with its mouth open. On the right, a larger white dog with brown ears and face is also jumping or running, looking towards the first dog. The background is a blurred landscape with some trees and a fence.

# Dog Breed Classification

...

Comparing transfer learning models

# The Stanford Dog Dataset

Contains 20k images from 120  
dog breeds, ruff-ly 150 images per  
breed



# Our Tasks

Transfer learning in PyTorch with  
ResNet & VGG

- Pre-process images in dataset
  - Create Trainer classes for pre-trained models
  - Fine tune models
  - Plot training and test accuracies between the models
  - Analyze results and compare the models
-

# Image Pre-processing

Original

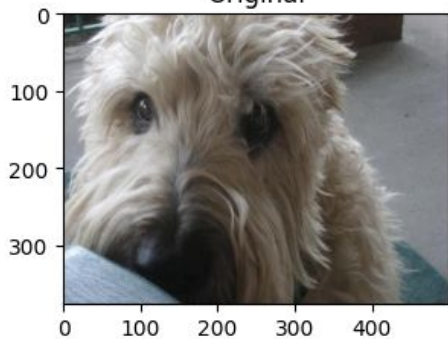
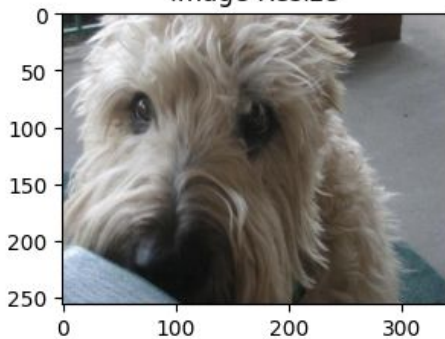
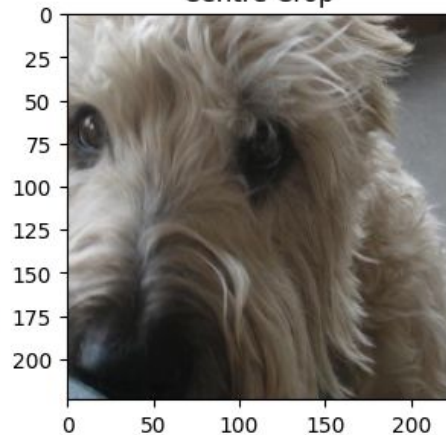


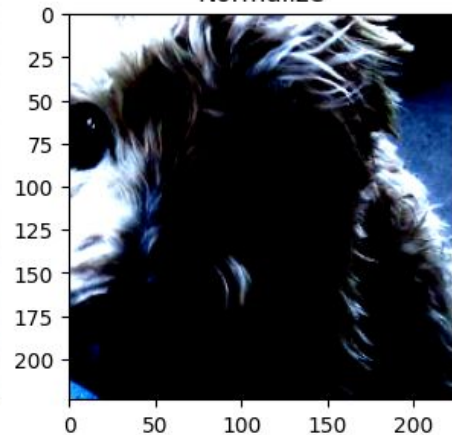
Image Resize



Centre Crop



Normalize



# Trainer Class for Pre-trained Models

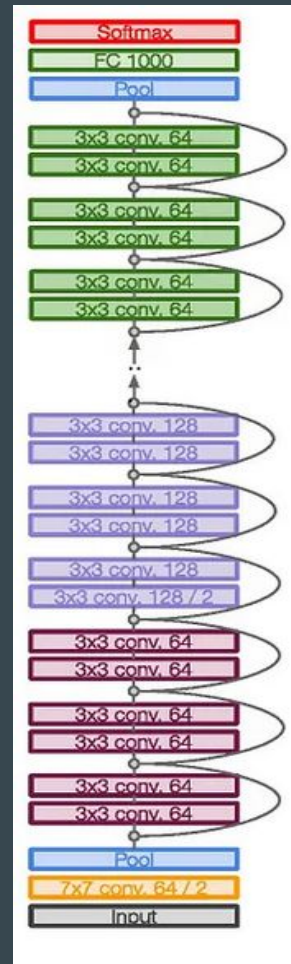
- Training and validation dataset is converted to dataloader
  - Contains images of size 224x224
- Cross entropy loss
- VGG and ResNet model is passed to respective class
  - ResNet-18, ResNet-50, VGG-16, VGG-11
  - Different fine-tuning methods per network
- Fully connected linear layer is changed for both networks
  - 120 output channels
- Train and save models

# Comparison of ResNet and VGG

- For our project we chose to do a comparative study of two pretrained machine learning models.
- The aim was to create a model which would take an image of a dog and classify it as the correct breed.
- The dataset we used was the Stanford Dogs dataset.

# The ResNet Model

- All ResNet models have a similar architecture, just with different layers.
- We used ResNet-50 and ResNet-18.
- The architecture consists of a number of residual blocks.
- Within each block we have multiple convolution layers, each of which are followed by batch normalization.
- All these blocks are stacked, so it becomes a very deep architecture.

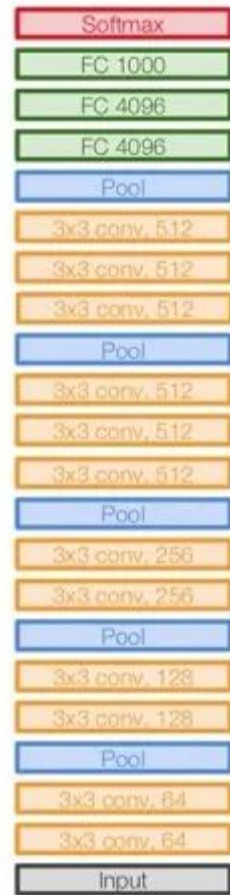






# The VGG model

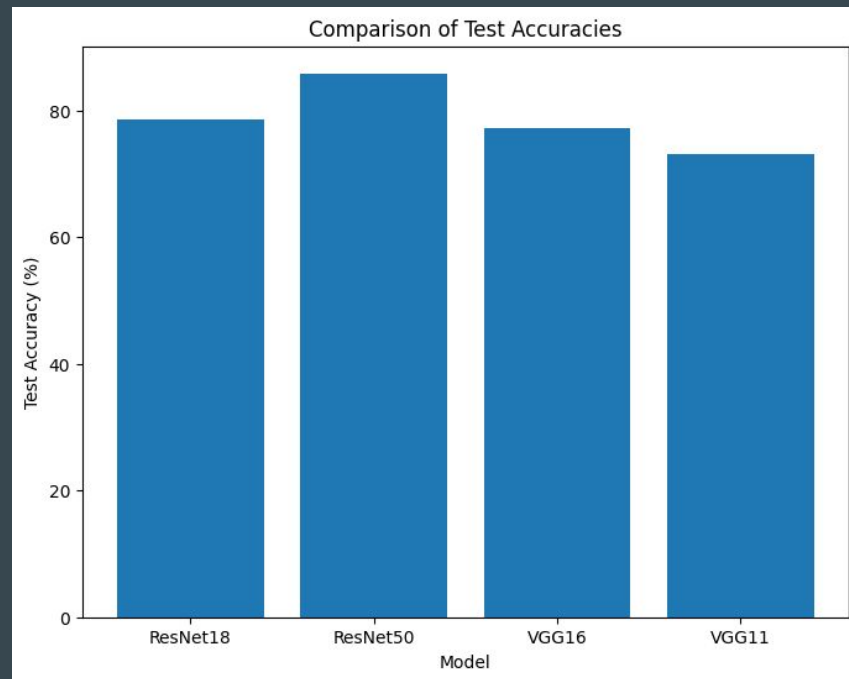
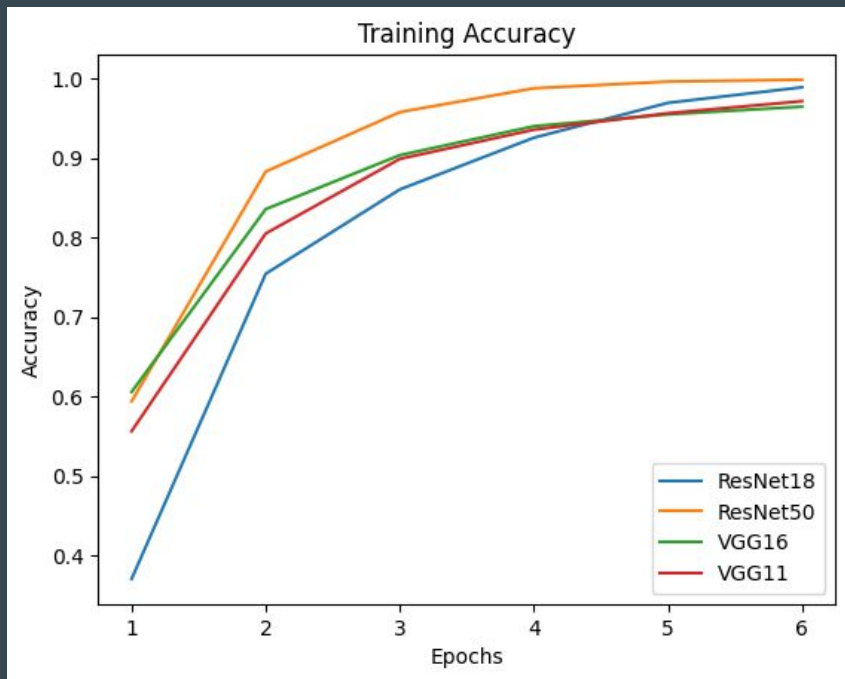
- All VGG models have a similar architecture, just with a different set of layers.
- We used VGG-11 and VGG-16.
- The VGG architecture uses small sets of convolution layers (2-4) followed by a pooling layer.
- This small grouping is stacked and followed by a fully connected layer.



VGG16



# Accuracy between models



# Training of ResNet models

```
# train model
```

```
resnet18_log = resnet18trainer.train(epochs)
```

```
[1 (116.7618s)]: train_loss=3.1977 train_acc=0.3707, val_loss=1.7912 val_acc=0.6651  
[2 (112.5373s)]: train_loss=1.2983 train_acc=0.7546, val_loss=1.1403 val_acc=0.7395  
[3 (110.0501s)]: train_loss=0.7498 train_acc=0.8606, val_loss=0.8885 val_acc=0.7715  
[4 (109.4918s)]: train_loss=0.4584 train_acc=0.9257, val_loss=0.8229 val_acc=0.7712  
[5 (110.6676s)]: train_loss=0.2780 train_acc=0.9696, val_loss=0.7666 val_acc=0.7858  
[6 (109.7251s)]: train_loss=0.1687 train_acc=0.9892, val_loss=0.7392 val_acc=0.7895  
== Total training time 669.2391 seconds ==
```

```
# train model
```

```
resnet50_log = resnet50trainer.train(epochs)
```

```
[1 (178.1292s)]: train_loss=2.1301 train_acc=0.5942, val_loss=0.7935 val_acc=0.8158  
[2 (179.3901s)]: train_loss=0.4763 train_acc=0.8831, val_loss=0.5481 val_acc=0.8496  
[3 (179.2760s)]: train_loss=0.1975 train_acc=0.9579, val_loss=0.4886 val_acc=0.8487  
[4 (178.7121s)]: train_loss=0.0819 train_acc=0.9879, val_loss=0.4601 val_acc=0.8583  
[5 (179.2338s)]: train_loss=0.0382 train_acc=0.9964, val_loss=0.4593 val_acc=0.8637  
[6 (178.5510s)]: train_loss=0.0204 train_acc=0.9988, val_loss=0.4662 val_acc=0.8618  
== Total training time 1073.2975 seconds ==
```

# Training of VGG models

```
# train model
```

```
vgg16_model = vgg16_model.to(device)
```

```
vgg16_trainer = VGG_Trainer(vgg16_model, train_dataset, val_dataset, 5e-5, 64)
```

```
vgg16_log = vgg16_trainer.train(epochs)
```

```
[1 (228.1726s)]: train_loss=1.4506 train_acc=0.6059, val_loss=0.6936 val_acc=0.7863  
[2 (227.6527s)]: train_loss=0.5095 train_acc=0.8358, val_loss=0.6828 val_acc=0.7867  
[3 (228.2055s)]: train_loss=0.2872 train_acc=0.9038, val_loss=0.7121 val_acc=0.7858  
[4 (228.2102s)]: train_loss=0.1825 train_acc=0.9402, val_loss=0.8456 val_acc=0.7585  
[5 (226.7153s)]: train_loss=0.1378 train_acc=0.9551, val_loss=0.8479 val_acc=0.7777  
[6 (228.2807s)]: train_loss=0.1013 train_acc=0.9648, val_loss=0.8041 val_acc=0.7799  
== Total training time 1367.2402 seconds ==
```

```
# train model
```

```
vgg11_model = vgg11_model.to(device)
```

```
vgg11_trainer = VGG_Trainer(vgg11_model, train_dataset, val_dataset, 5e-5, 64)
```

```
vgg11_log = vgg11_trainer.train(epochs)
```

```
[1 (156.2669s)]: train_loss=1.6267 train_acc=0.5566, val_loss=0.8424 val_acc=0.7508  
[2 (156.7959s)]: train_loss=0.6028 train_acc=0.8052, val_loss=0.7919 val_acc=0.7607  
[3 (156.1768s)]: train_loss=0.2997 train_acc=0.8991, val_loss=0.7921 val_acc=0.7595  
[4 (155.6686s)]: train_loss=0.1827 train_acc=0.9358, val_loss=0.9015 val_acc=0.7375  
[5 (156.2055s)]: train_loss=0.1320 train_acc=0.9566, val_loss=0.9846 val_acc=0.7386  
[6 (156.2990s)]: train_loss=0.0866 train_acc=0.9717, val_loss=1.0465 val_acc=0.7342  
== Total training time 937.4153 seconds ==
```

# Deployment

- Python file that identifies the breed of dog
  - Dog images would be taken from a directory
  - Each image will be processed before evaluation
- Loads saved model from file
- Images are classified and the breed is named
- Possible future integration to front-end apps or to other models

# References

- <https://cv-tricks.com/keras/understand-implement-resnets/>
- <https://towardsdatascience.com/architecture-comparison-of-alexnet-vggnet-resnet-inception-densenet-beb8b116866d>
- <http://vision.stanford.edu/aditya86/ImageNetDogs/>